

Towards Certified Program Obfuscation

Weiyun Lu

*School of Electrical Engineering and Computer Science
University of Ottawa
Ottawa, Canada
WLU058@uottawa.ca*

Bahman Sistany

*Cloakware Research
Irdeto Canada
Ottawa, Canada
bahman.sistany@irdeto.com*

Amy Felty

*School of Electrical Engineering and Computer Science
University of Ottawa
Ottawa, Canada
afelty@uottawa.ca*

Philip Scott

*School of Electrical Engineering and Computer Science
University of Ottawa
Ottawa, Canada
philip.scott@uottawa.ca*

Abstract—How can one perform code transformations such as obfuscating transformations or optimizing transformations on code that is assumed to be correct with respect to certain specified behaviour? Will the transformed code preserve the specified behaviour as one expects?

To achieve the highest levels of assurance that the transformations have maintained correctness of the code, is to prove the two versions of the program (before and after the transformation) is equivalent. Total equivalency between the two versions of a program certainly implies the correctness of any specified properties of interest but what if we directly try to show validity of these properties on the transformed program?

In this thesis, we lay the foundation to study and reason about code obfuscating transformations and show how the preservation of certain behaviours may be “certified”. To this end, we apply techniques of formal specification and verification, by using the Coq Proof Assistant and IMP (a simple imperative language within it), to formulate what it means for a program’s semantics to be preserved by an obfuscating transformation, and give formal machine-checked proofs that these properties hold.

We describe our work on opaque predicates, a simple control flow obfuscation, and elements of control flow flattening transformation, a more complex control flow obfuscation. Along the way, we employ Hoare logic as our foundational specification language, as well as augment the IMP language with Switch statements. We also define a lower-level flowchart language to wrap around IMP for modelling certain flattening transformations, treating blocks of codes as objects in their own right.

Index Terms—obfuscation, verification, security, correctness, Coq, proof

1. Introduction

1.1. Background and Motivation

We expect our systems including software systems to function “correctly”. By “correctly”, we mean that a

system will behave according to explicit and/or implicit expectations or said another way to its written and/or unwritten specifications. Typically, extensive testing is done to increase the confidence in correct functionality of a piece of software but alas testing is known to be based on inductive reasoning where more tests passing can only increase the likelihood of correctness, so positive testing results are not a proof of correctness. In systems where more assurance of correctness is required various types of deductive reasoning is used. These are formal verification methods based on theoretical foundations rooted in logic. It is important to note that, formal verification transfers the problem of confidence in program correctness to the problem of confidence in specification correctness, so it is not a silver bullet however since specifications are often smaller and less complex to express, we have successfully reduced the trusted computing base (TCB) and increased our chances of achieving correctness.

Formal verification based methods, used to show a (software) system behaves as its specification says, typically employ a specification language based on the familiar “assertions”. A specification is typically expressed in some variation of first order logic and the verification system will deductively try to prove the assertions correct or signal that they don’t hold. This is a rather elaborate process where assertions (general propositional statements about program fragments that are expected to hold) are used to generate verification conditions (VC), logic formulas, that are then fed into a satisfiability modulo theories (SMT) solver, either behind the scenes in a verification backend or in more visible to the verification expert. VC generation for program verification goes back to at least Hoare’s triples, Eiffel style contracts and proof-carrying-code (PCC) of Necula [9].

How can one perform code transformations such as obfuscating transformations or optimizing transformations on code that is assumed to be correct with respect to certain specified behaviour (expressed in some assertion language) while preserving the correctness of the specified behaviour?

To achieve the highest levels of assurance that the transformations have maintained correctness of the code,

is to prove the two versions of the program (before and after the transformation) is equivalent. Proving equivalency is doable in certain cases but in general is still extremely hard to do for general programs. Despite the difficulty, there are now formally verified compilers such as CompCert. However, the problem with verifying realistic systems such as a compiler is scale. CompCert verification of semantic equivalence between C and generated assembly took several man years to complete. The cost of using formal verification for mere mortals (on realistic systems) is still high.

Total equivalency between the two versions of a program certainly implies the correctness of any program properties of interest but what if we directly try to show validity of these properties? What if we limit ourselves to only proving properties of interest in the “before” version of a program are maintained in the “after” version of the program (after a transformation is applied)?

Certain simple transformations simply don’t invalidate expressed properties about the “before” version versus the “after” version. Below the program snippet in listing [1] asserts that $y > 2$ which we can verify visually to be true. In the snippet in listing [2] we use a simple obfuscating transformation called variable splitting where we have split the variable x into two other variables $x1$ and $x2$ and we see that (visually) the assertion $y > 2$ still holds.

```
x = 2; y = 5;
y = x + y;
assert(y > 2);
```

Listing 1: Original Code

```
x1 = 1; x2 = 1; y = 5;
y = x1 + y; y = x2 + y;
assert(y > 2);
```

Listing 2: Obfuscated Code

In general, though, most transformations, whether optimizations or obfuscations but specially obfuscations, invalidate assertions that hold true about the “before” version. Obfuscation is especially troublesome because the goal of obfuscation is to hide the functionality of the code from prying eyes while maintaining the functionality of the “before” program. “Prying eyes” could as easily be the same as some kind of static analysis tool where an attacker is trying to determine certain facts about the code and obfuscation is trying to make this difficult. The program in listing [3] is correct with respect to the assertion that is expressed (e.g. $z == 30$) as is evident by simple inspection of the code. The program snippet in listing [4] is the “after” program where a non-linear opaque predicate transformation has been applied to hide the fact that at program’s end, value of z is in fact 30. We can see that the transformation makes it a bit harder to see that the assertion still holds but knowing the fact that $\forall x \in \mathbb{Z}, ((x^2 + x) \bmod 2) == 0$, we can deduce that the assertion does hold and the value of z is in fact still 30.

This paper describes steps towards implementing a framework in the Coq Proof Assistant and based on IMP [12], a simple imperative language. to study obfuscating

transformations, their impact on programs and how specified behaviour may be preserved beyond the transformations. A of number of initial goals and principles drove the direction of this research: 1) Not reinventing the wheel: start out with IMP a familiar small imperative language implemented in Coq and use its accompanying formalized semantics in [12]. 2) Accessibility to as wide an audience as possible: an obvious option was to use CompCert and Clight as [2] have done. We would have started with lots of proofs and formalisms for free (already done by the CompCert team) however the significant learning curve associated with learning CompCert infrastructure seemed prohibitive. We deemed IMP and Coq much more accessible. 3) Extendibility of the framework: Following the lead of [12] where a number of extensions to IMP are easily added and studied, we wanted the ability to build our obfuscation infrastructure incrementally on top of IMP.

Keeping these research goals in mind the contributions of this paper are the following:

- We consider different formulations of what it means for a transformation to be semantics-preserving, including complete state equivalence as well as Hoare logic equivalence. In this particular setting, the latter is a novel approach, and we give examples of its use with opaque predicate transformations. In addition, use of Hoare logic in this context leads to establishing our main approach to “certifying obfuscating transformations”: our obfuscating transformations will be “decorated” à la Pierce’s [12] with additional assertions and their proofs.
- We give clear and detailed explanations of the proofs and tactics in Coq, which, to the best of our knowledge, the existing literature does not, thus providing an accessible explanation of not just obfuscation techniques, but also in tandem with its formalization and verification inside Coq. This follows from research goal 2.
- We begin with a minimal imperative programming language inside Coq for reasoning about programs and their transformations, and then augment it as needed for control flow flattening algorithms, first by augmenting its syntax and semantics with switch statements, and then by defining a lower-level flowchart language that wraps around blocks of code in order to model real-world intermediate languages used in obfuscation tools. This follows our research goals 2 and 3.

```
int main (int argc, char *argv[])
{
    unsigned int x = 10;
    unsigned int y = 20;
    unsigned int z = 0;

    z = x + y;
    assert(z == 30);
    return 0;
}
```

Listing 3: Original Code

```

int main (int argc, char *argv[])
{
    unsigned int x = 10;
    unsigned int y = 20;
    unsigned int z = 0;

    unsigned int a = ((unsigned int)argc);
    unsigned int w = a * a;

    w = a + w;
    w = w % 2;
    if (w == 0)
    {
        z = x + y;
    }
    else
    {
        z = y - x;
    }
    assert(z == 30);
    return 0;
}

```

Listing 4: Obfuscated Code

2. Background — formal verification

In this chapter we provide necessary background information used for formal verification — namely, the Coq proof assistant, the simple imperative language IMP defined inside Coq, and Hoare logic for reasoning about pre- and post-conditions of programs.

2.1. The Coq proof assistant

3. Opaque predicates in IMP/Coq

We now enter the main topic of the thesis proper, formalizing and certifying the opaque predicate transformation mentioned earlier.

3.1. Opaque predicate obfuscation

An *opaque predicate* [13] is a predicate¹ that always evaluates to either *true* or *false* and the truth-value of which is known to the transformation but hard to deduce by an attacker [14]. The code under the *false* branch is never evaluated at runtime so *opaque predicates* incur no runtime performance.

Of course, the absolutely most basic opaque predicates are just the boolean constants *true* and *false* themselves, but these are not very useful in practice because it is immediately obvious what is happening in the program, and neither the simplest of humans nor tools will be fooled. For a more advanced treatment of opaque predicates and how they may be broken see [14].

1. This could be any statement in a program that could evaluate to true or false, but we will only be concerned with arithmetic formulas in this thesis.

An *opaque predicate* transformation takes as inputs a program to be obfuscated, c_1 , an opaque predicate P , and a dummy program² c_2 , and returns the program

$$\text{IFB } (P \ x) \text{ THEN } c_1 \text{ ELSE } c_2 \text{ FI.}$$

In the Section 3, we describe our initial (straightforward, naive) attempt, explicitly defining the transformation to introduce the lines of code that assign variables associated with the opaque predicate (as one may naturally expect to write code in a typical imperative language), and see that trying to state a general theorem about command equivalence ends up being problematic.

However, we then discuss how this spawned two ideas in different directions that rectify the issue. On the one hand, we use Hoare logic with this first formulation, in Section ??, to prove weaker conditions of a transformation than total command equivalence. On the other hand, in Section ?? we reformulate the transformation to rely on values already existing in the state of the program, with the view that one may be applying an opaque predicate transformation to a small piece of code somewhere within a much larger program that would have such values floating around in the state already.

Finally in Section ??, we again employ Hoare logic to give a formal example of how an attacker who does not know about the opaque predicate’s constant truth valuation, but otherwise can analyze (using static analysis) the program, ends up gaining weaker knowledge because of it.

All code for this chapter is in the file `OBFS_opaque_predicate.v` [?].

4. Command equivalence

4.1. Factorial program (countdown nonzero formulation)

The following IMP program computes the factorial of a nonzero natural number. The input is read from X , temporary values are stored as Z , and the factorial of the input is stored in Y as the output.

```

Definition fact_nonzero : com :=
  Z ::= X;;
  Y ::= 1;;
  WHILE ! (Z <= 1) DO
    Y ::= Y * Z;;
    Z ::= Z - 1
  END.

```

The choice of factorial program as a candidate for examples of obfuscation is somewhat arbitrary. It works well for illustrative purposes, however, as it is neither too complex nor completely trivial.

4.2. Example

The *fact_nonzero* program with input $X = 3$ yields output $Y = 6$. However, the story is not quite so simple (it

2. It’s not known to an attacker, a priori, that it’s a dummy program. In practice, c_2 should be constructed so that it is not obvious; e.g. c_2 should not be simply an empty program, but should be complicated enough that it looks like it could feasibly be intended to be executed.

is true that input $X = 3$ yields $Y = 6$, but as one can see in the Coq example, the state keeps track of the value of every variable involved in the program.). The specification of this statement in Coq is

```
Example factorial_3:
fact_nonzero / { X --> 3 }
{ X --> 3; Z --> 3; Y --> 1; Y --> 3;
  Z --> 2; Y --> 6; Z --> 1 }.
```

Note that formally, the final state holds the information of every intermediate assignment made by the program. We can discern the output $Y = 6$ by the fact that this is the rightmost case of a value being assigned to Y . But wait, there's more! We said earlier that in Coq, an example is no different from a proposition or a theorem in anything but name, so we must actually give a proof³. Moreover, since command evaluation is relational and not functional (recall the reason for this is the possibility of non-terminating While loops), we must build the proof out step by step (see [10] for proofs).

For this section, we'll use as a simple opaque predicate namely,

$$\forall x. (x * x + x + x + 1) = (x + 1) * (x + 1).$$

We now define an opaque predicate transformation with our running example. For the purposes of making the proofs easier to work with, and also to add a slight additional touch of obfuscation, we split up these assignments over multiple lines, as follows.

```
Definition opaque_trans x c1 c2 :=
  X' ::= (ANum x) ;;
  Z' ::= X' * X' ;;
  Z' ::= Z' + X' ;;
  Z' ::= Z' + X' ;;
  Z' ::= Z' + 1 ;;
  Z'' ::= X' + 1 ;;
  Z'' ::= Z'' * Z'' ;;
  IFB (BEq Z' Z'') THEN c1 ELSE c2 FI.
```

holds and executes c_2 otherwise. Of course, the above is true for all x , so the resulting program should be the same as c_1 . We'd like to claim that a program transformed by `opaque_trans` is equivalent to the original.

The observant logically inclined reader should, at this point, now be suspicious about taking this claim at face value. What do we mean when we say the transformed program should be “the same”? The next example, which shows what happens when `opaque_trans` is applied to the `factorial_3` example, elucidates the necessity to be precise. First, however, we will need a few lemmas that show our opaque predicate is indeed such, in various incarnations to be used in proofs.

Lemma 4.1. It is indeed the case that

$$\forall x \in \mathbb{N}, (x * x + x + x + 1) = (x + 1) * (x + 1).$$

3. This really is an example, to us. But just because one declares “here is an example of X ” does not mean that X is necessarily true. In Coq, a proof must still be constructed.

For example, in natural language, one can say “An example of a prime number is 20051”. But this isn't immediately obvious, and one still needs to prove that example, for instance, by writing a program that tries to divide it by every number up to its square root.

4.3. Example

For any $x \in \mathbb{N}$ and any program c_2 , `opaque_trans x fact_nonzero c2` with input $X = 3$ yields output $Y = 6$. In Coq, however, it looks as follows.

```
Example factorial\_3\_opaque\_trans:
forall x c2,
opaque_trans x fact_nonzero
c2 / { X --> 3 }
{ X --> 3; X' --> x; Z' --> x * x;
  Z' --> x * x + x;
  Z' --> x * x + x + x;
  Z' --> x * x + x + x + 1;
  Z'' --> x + 1;
  Z'' --> (x + 1) * (x + 1);
  Z --> 3; Y --> 1; Y --> 3;
  Z --> 2; Y --> 6; Z --> 1 }.
```

After instantiating variables and unfolding the definitions, the transformed program looks like:

```
X' ::= x;;
Z' ::= X' * X';;
Z' ::= Z' + X';;
Z' ::= Z' + X';;
Z' ::= Z' + 1;;
Z'' ::= X' + 1;;
Z'' ::= Z'' * Z'';;
IFB Z' = Z''
THEN Z ::= X;;
  Y ::= 1;;
  WHILE ! (Z <= 1) DO Y ::= Y * Z;;
  Z ::= Z - 1
END
ELSE c2 FI
```

We cannot use *cequiv* (??) — that is, we can't use it with the current formulation of the transformation) — since new variables and assignments are introduced and kept track of in the definition of the state, even if we ultimately don't care about them.

Thus we were not ultimately successful, in this initial approach, in formulating a statement with command equivalence (??). We'll revisit this in Section ??.

5. Related works

There have been three papers, in all of which Sandrine Blazy (Université de Rennes 1) appears as a coauthor, that study code obfuscation in Coq.

Towards a formally verified obfuscating compiler

Towards a formally verified obfuscating compiler [1] also uses IMP as the language for obfuscation, but studies data obfuscation techniques, as opposed to the control obfuscation techniques which opaque predicates and control flow flattening fall under [11].

The first particular transformation studied herein is obfuscating integer constants, wherein all integer values are substituted by different ones in a distorted semantics

using an obfuscating function $O : \mathbb{N} \rightarrow \mathbb{N}$. The other discussed is variable encoding, which changes the names of variables. A real-life application of this could be, for instance, to change a descriptive variable name like `account_balance` to a string of gibberish.

This is an inherently different class of techniques from the ones studied in the present work, and one can make a simple combinatorial argument that putting them together in the same obfuscation transformation would generate a synergistic effect on making a program more difficult to analyze.

Formal verification of control-flow graph flattening

Formal verification of control-flow graph flattening [2] also studies control flow flattening, but the authors use the Clight language of CompCert (the formally verified C compiler in Coq, discussed in Chapter Clight is the first intermediate language in the CompCert compiler, and the strategy used was to prove the correctness of the obfuscation strictly there, from which CompCert’s own proofs of semantic preservation give the correctness of the rest of the compilation process “for free”).

On the one hand, this makes the work less elementary and less accessible, as it works with a nontrivial subset of the real C language, but on the other it is clear evidence that formal verification of obfuscation techniques need not be restricted to a small language like IMP (which would never be used in real software development), and other real-world practicalities considered in this paper include simulation techniques and analysis of running time.

This work also discusses some techniques for combining obfuscation techniques, such as splitting a switching variable into two different variables that are updated at different points of a program, as well as randomly encoding the values of the switch cases so that they are not just consecutive numbers beginning with 1. These are necessary considerations, since we need to think one level higher about our attackers, and obfuscate the fact that we are obfuscating particular parts of our code with CFG flattening in the first place!

In comparing this work to ours, the present author believes there is merit both in the IMP and the CompCert routes. In the former, the language used is of minimal complexity, which allows not only for specifications and proofs of transformations to be developed quicker without being bogged down in unnecessarily complicated features of the underlying language, but is also better suited for pedagogical purposes. IMP is also Turing complete, so from a theoretical point of view there is no loss of generality in proofs made using it — they can always be adapted to CompCert later. But on the other hand, CompCert is, of course, closer to languages that would be of interest to real-world software development and so more practical in that sense.

The authors ran into a similar issue as in the present work of needing to separate switching variables from those in the program to be transformed, but their solution was different — they instead use a function to parse the program to be transformed and generate a fresh variable which doesn’t appear there to be used for the transformation. From a practical point of view, this is perhaps

more natural, and in line with how a real obfuscating tool would function — generating new variables rather than demand that a certain specifically named variable doesn’t exist in the source program. Theoretically, though, these are equivalent, since any program can contain only finitely many variable names, and there are an infinite number to choose from.

Formal verification of a program obfuscation based on mixed boolean-arithmetic expressions

Formal verification of a program obfuscation based on mixed boolean-arithmetic expressions [3] mixed boolean-arithmetic continues to work in Clight, which studies obfuscations that involve mixing arithmetic operators and bitwise boolean operators. This is another data obfuscation which appears frequently in real-world binary code, but as it is based on features wildly beyond the capabilities of IMP, a detailed discussion is beyond the scope of the present work.

References

- [1] Sandrine Blazy and Roberto Giacobazzi. Towards a formally verified obfuscating compiler. In 2nd ACM SIGPLAN Software security and protection workshop, SPP 2012, St. Petersburg, FL, USA, 2012. HAL (<https://hal.archives-ouvertes.fr/>).
- [2] Sandrine Blazy and Alix Trieu. Formal verification of control-flow graph flattening. In Proceedings of the 5th ACM SIGPLAN Conference on Certified Programs and Proofs, CPP 2016, pages 176–187, New York, NY, USA, 2016. ACM.
- [3] Sandrine Blazy and Remi Hutin. Formal verification of a program obfuscation based on mixed boolean-arithmetic expressions. In Proceedings of the 8th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2019, Cascais, Portugal, January 14–15, 2019, pages 196–208, 2019.
- [4] Xavier Leroy. Formal verification of a realistic compiler. *Commun. ACM* 52(7): 107–115 (2009)
- [5] I. S. Jacobs and C. P. Bean, “Fine particles, thin films and exchange anisotropy,” in *Magnetism*, vol. III, G. T. Rado and H. Suhl, Eds. New York: Academic, 1963, pp. 271–350.
- [6] K. Elissa, “Title of paper if known,” unpublished.
- [7] R. Nicole, “Title of paper with only first word capitalized,” *J. Name Stand. Abbrev.*, in press.
- [8] Y. Yorozu, M. Hirano, K. Oka, and Y. Tagawa, “Electron spectroscopy studies on magneto-optical media and plastic substrate interface,” *IEEE Transl. J. Magn. Japan*, vol. 2, pp. 740–741, August 1987 [Digests 9th Annual Conf. Magnetism Japan, p. 301, 1982].
- [9] George C. Necula: Proof-Carrying Code. *POPL* 1997: 106–119.
- [10] Weiyun Lu, “GitHub repository to accompany the present thesis,” <https://github.com/weiyunlu/coq-certified-obfuscation>, 2019
- [11] C. Collberg, C. Thomborson, and D. Low. A Taxonomy of Obfuscating Transformations. Technical Report 148, Department of Computer Science, University of Auckland, July 1997.
- [12] Benjamin C. Pierce, Arthur Azevedo de Amorim, Chris Casinghino, Marco Gaboardi, Michael Greenberg, Cat? alin Hri?cu, Vilhelm Sjberg, Andrew ? Tolmach, and Brent Yorgey. *Programming Language Foundations*. Software Foundations series, volume 2. Electronic textbook, May 2018.
- [13] Christian S. Collberg, Jasvir Nagra: Surreptitious Software - Obfuscation, Watermarking, and Tamperproofing for Software Protection. Addison-Wesley Software Security Series, Addison-Wesley 2010, ISBN 978-0-321-54925-9, pp. I–XXVII, 1–748
- [14] Mila Dalla Preda, Matias Madou, Koen De Bosschere, Roberto Giacobazzi: Opaque Predicates Detection by Abstract Interpretation. *AMAST* 2006: 81–95